

NAME : DIVYA AC

USN : 1SV21CS029

TEAM 09

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from matplotlib import pyplot as plt

data = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')

df = pd.DataFrame(data)

print(df.columns)

Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm',
       'PetalWidthCm',
       'Species'],
      dtype='object')

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

data = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv', sep=';')

df = df.reset_index()

print(df.columns)

Index(['index', 'Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm',
       'PetalWidthCm', 'Species'],
      dtype='object')

import pandas as pd

# Assuming your CSV file has ';' as the delimiter
data = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv', sep=';')

# Print out the column names to verify
print(data.columns)

Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm', 'Species'],
      dtype='object')

# Mount Google Drive if you're using Google Colab
from google.colab import drive
```

```
drive.mount('/content/drive')
```

```
import pandas as pd
```

```
# Path to your CSV file in Google Drive
```

```
csv_file_path = '/content/drive/MyDrive/archive (10)/iris.csv'
```

```
# Reading the CSV file; adjust the separator based on your file
```

```
data = pd.read_csv(csv_file_path, sep=',')
```

```
# Verify the column names and data
```

```
print(data.head())
```

```
# Select only two specific columns for independent variables (X)
```

```
X = data[['SepalLengthCm', 'SepalWidthCm']] # Independent variables  
(features)
```

```
# Dependent variable (target)
```

```
y = data['Species'] # Assuming 'Species' is the target variable
```

```
# Print first few rows of X and y to verify
```

```
print("Features (X):")
```

```
print(X.head())
```

```
print("\nTarget variable (y):")
```

```
print(y.head())
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

Features (X):

	SepalLengthCm	SepalWidthCm
0	5.1	3.5
1	4.9	3.0
2	4.7	3.2
3	4.6	3.1
4	5.0	3.6

Target variable (y):

0	Iris-setosa
1	Iris-setosa
2	Iris-setosa
3	Iris-setosa

```
4    Iris-setosa
Name: Species, dtype: object
```

```
k=3
knn=KNeighborsClassifier(n_neighbors=k)
knn.fit(X,y)
```

```
KNeighborsClassifier(n_neighbors=3)
```

```
new_data = np.array([[5.1,3.5]])
prediction = knn.predict(new_data)
if prediction[0] == 1:
    print("Iris-setosa")
else:
    print("Iris-versicolour")
```

```
Iris-versicolour
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X
does not have valid feature names, but KNeighborsClassifier was fitted with
feature names
  warnings.warn(
```

```
# Load the dataset
```

```
file_path = ('/content/drive/MyDrive/archive (10)/iris.csv')
```

```
# Update this path accordingly
```

```
df = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')
```

```
# Define independent variables (features) and dependent variable
```

```
X = df[['SepalLengthCm', 'SepalWidthCm']]
```

```
y = df['Species'] # Assuming 'Species' is the target variable
```

```
# Encode the target variable 'Species' into numerical values
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
y = le.fit_transform(y)
```

```
# Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Create the Linear Regression model (for demonstration, though not ideal for
this problem)
```

```
model = LinearRegression()
```

```
# Train the model
```

```
model.fit(X_train, y_train)
```

```
LinearRegression()
```

```
# Make predictions on the testing set
```

```
y_pred = model.predict(X_test)
```

Evaluate the model (these metrics might not be appropriate for categorical predictions)

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

For classification problems, consider using Logistic Regression, Decision Trees, or other suitable algorithms.

```
Mean Squared Error: 0.16908805917847763
R-squared: 0.7580616005395391
```

```
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

```
Coefficients: [ 0.72039588 -0.66538649]
Intercept: -1.1588138946175666
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score
# Load dataset
df = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')
# Define independent variables (features) and dependent variable
# (target) # This line seems to be an unintended command, commenting it out
X = df[['SepalLengthCm', 'SepalWidthCm']]
y = df['Species']
```

Encode the target variable 'Species' into numerical values

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y) # Convert target to numerical
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

Initialize k-NN regressor (set k=3 for example)

```
k = 3
```

```
knn_regressor = KNeighborsRegressor(n_neighbors=k)
```

Train the model

```
knn_regressor.fit(X_train_scaled, y_train)
```

Predict on the test set

```
y_pred = knn_regressor.predict(X_test_scaled) # Now y_pred will be numerical
```

Evaluate performance (e.g., using RMSE and R^2)

```
rmse = mean_squared_error(y_test, y_pred, squared=False)
r2 = r2_score(y_test, y_pred)
print(rmse)
print(r2)
```

```
0.3751542892474251
```

```
0.798622151563328
```

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score
```

```
# Load dataset
```

```
df =pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')
```

```
# Display the first few rows to understand the structure
```

```
print(df.head())
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
X = df[['SepalLengthCm', 'SepalWidthCm']]
```

```
y = df['Species']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize logistic regression model
```

```
log_reg = LogisticRegression(random_state=42)
```

```
# Train the model
```

```
log_reg.fit(X_train_scaled, y_train)
```

```
LogisticRegression(random_state=42)
```

```
# Predict on the test set
```

```
y_pred = log_reg.predict(X_test_scaled)
```

```
# Evaluate performance
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
class_report = classification_report(y_test, y_pred)
```

```
print(f'Accuracy: {accuracy}')
```

```
print(f'Confusion Matrix:\n{conf_matrix}')
```

```
print(f'Classification Report:\n{class_report}')
```

Accuracy: 0.9

Confusion Matrix:

```
[[10  0  0]
 [ 0  7  2]
 [ 0  1 10]]
```

Classification Report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	0.88	0.78	0.82	9
Iris-virginica	0.83	0.91	0.87	11
accuracy			0.90	30
macro avg	0.90	0.90	0.90	30
weighted avg	0.90	0.90	0.90	30

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
# Load dataset
df = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')
# Display the first few rows to understand the structure
print(df.head())
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

... (rest of your code)

```
X = df[['SepalLengthCm', 'SepalWidthCm']]
y = df['Species']
```

```
# Encode the target variable 'Species' into numerical values
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y) # Convert target to numerical
```

```
# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Initialize decision tree regressor
dt_regressor = DecisionTreeRegressor(random_state=42)
```

```

# Train the model
dt_regressor.fit(X_train, y_train)
# ... (rest of your code)

DecisionTreeRegressor(random_state=42)

y_pred = dt_regressor.predict(X_test)
# Evaluate performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error (MSE): {mse}')
print(f'R-squared (R2): {r2}')

Mean Squared Error (MSE): 0.29814814814814816
R-squared (R2): 0.5733969263381027

```

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
# Load dataset
df = pd.read_csv('/content/drive/MyDrive/archive (10)/iris.csv')
# Display the first few rows to understand the structure
print(df.head())

```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```

# ... (rest of your code)

X = df[['SepalLengthCm', 'SepalWidthCm']]
y = df['Species']

# Encode the target variable 'Species' into numerical values
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y) # Convert target to numerical

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Initialize decision tree regressor
rf_regressor = RandomForestRegressor(random_state=42)

```

```
# Train the model
rf_regressor.fit(X_train, y_train)
# ... (rest of your code)

RandomForestRegressor(random_state=42)

y_pred = rf_regressor.predict(X_test)
# Evaluate performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error (MSE): {mse}')
print(f'R-squared (R2): {r2}')
```

Mean Squared Error (MSE): 0.192646804138322
R-squared (R2): 0.7243527444761688